

problem: determine what lines to investigate.

guideline: investigate moves that are most likely to be played.

concern: only broad search.
remedy: approximate minimax more closely
concern: investigates known lines
remedy: use uncertainty... somehow

Objective:

pick s : ~~$\arg\max_s P(s \text{ played}) \Delta \text{uncertainty}(s)$~~

to minimize $\sum_{s' \neq s} P(s') \text{uncertainty}(s')$

need:

- access to P
- prediction of $\Delta \text{uncertainty}$
- fast propagation of $\Delta \text{uncertainty}$ benefit
- back propagation

access to P

occasionally update whole tree

if frequency = $O(\log |S|)$, no asymptotic decrease in efficiency

try frequency = $O(\sqrt{|S|})$

overhead = $\sum_{i=1}^{\sqrt{|S|}} i^2 = O(|S|^{1.5})$ unacceptable

clean_ep $\leftarrow 1$

upon clean: clean_ep \leftarrow clean_ep $\cdot 1.1 + 1$

prediction of Δ uncertainty:
initial = max uncertainty (always)
end has to be predicted
use all uncertainties to train,
weight "nearby" samples asymptotically more
"nearby" = at same depth?

OR assume Δ uncertainty constant across all
states

backprop of uncertainty loss

Uncertainty (parent) upon observation of child

- re

